

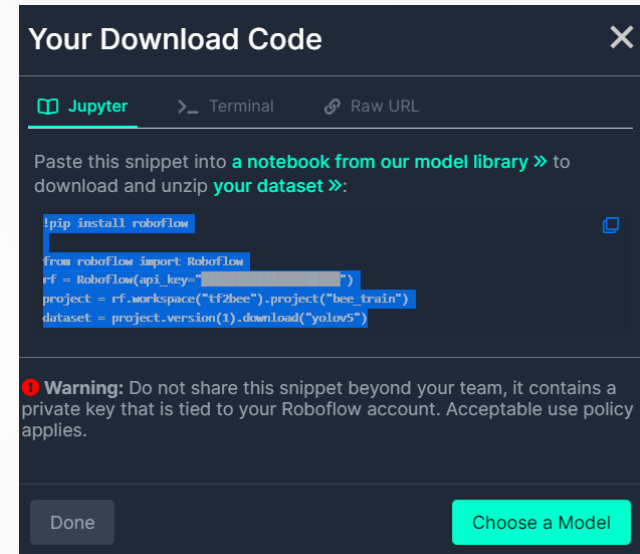
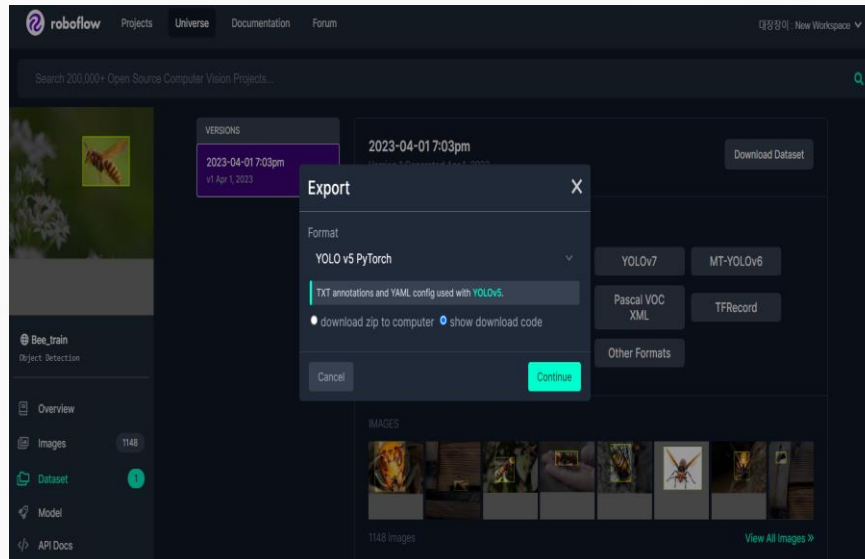
수집된 이미지를 바탕으로 객체인식 실습

이 준 혁 (2023250419)
산 업 인 공 지 능 학 과

진행순서

1. Datasets 가져오기
2. Yolo 클론 설치 및 프로그램 패키지 설치
3. 이미지 저장 리스트생성 및 이미지 개수확인
4. 트레이닝셋과 밸리데이션 셋 나누기
5. 위치 변경
6. 트레이닝 시키기
7. 트레이닝 완료하고 확인하기
8. 이미지 넣고 추론하기
9. 내가 원하는 이미지를 넣고 확인하기

1. 데이터셋 가져오기 (1)



로보플로우사이트에서 데이터셋 구하기

1. 데이터셋 가져오기 (2)

[5] #말발에 데이터셋을 가져오기 위한 코드(roboflow에서 dataset을 제공하고 있음)
!pip install roboflow

```
from roboflow import Roboflow  
rf = Roboflow(api_key="fNmB1Zp1Cdu4#0EN0rmQ")  
project = rf.workspace("tf2bee").project("bee_train")  
dataset = project.version(1).download("yolov5")
```

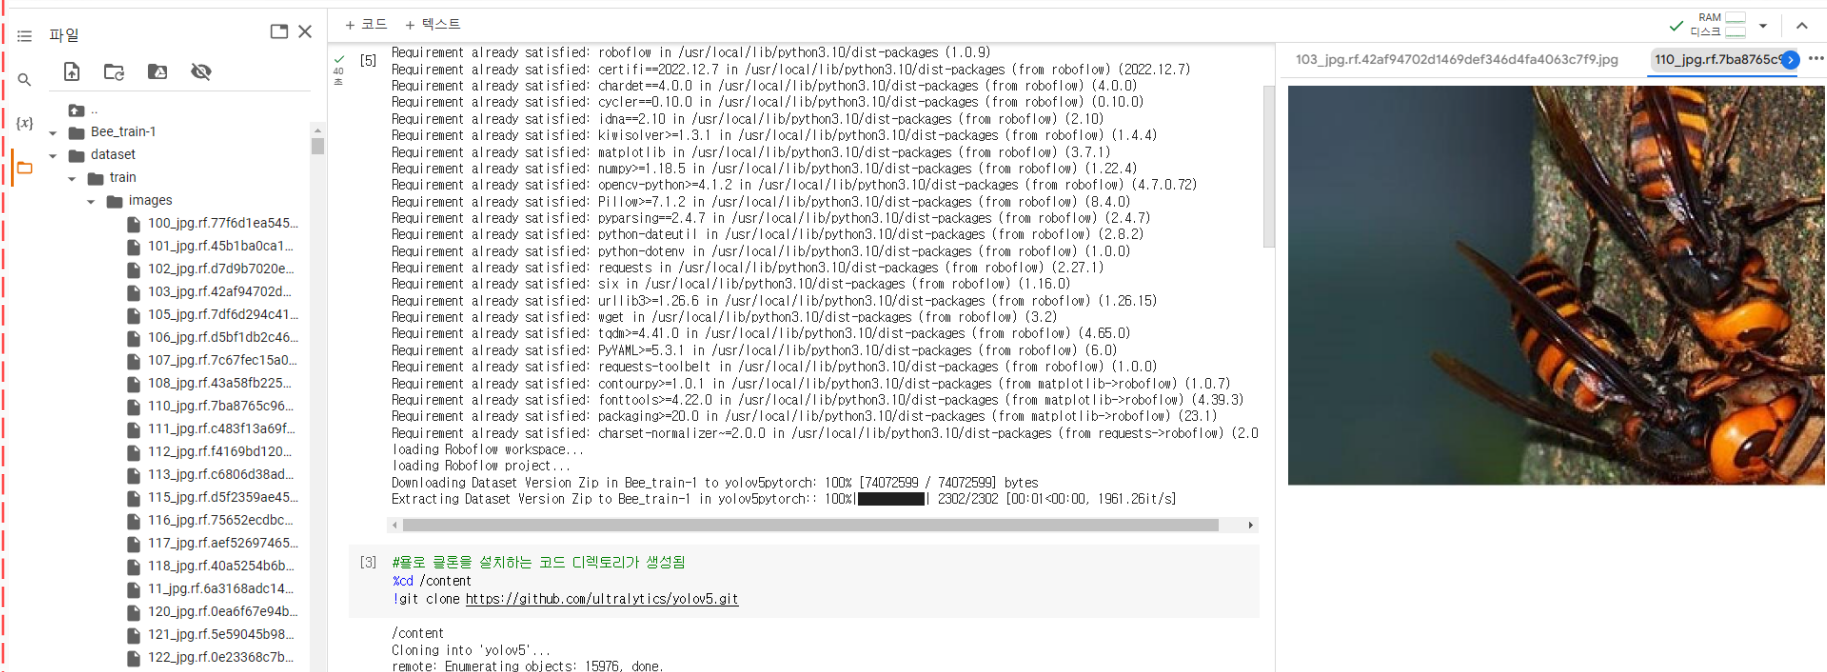
Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>
Requirement already satisfied: roboflow in /usr/local/lib/python3.10/dist-packages (1.0.9)
Requirement already satisfied: certifi==2022.12.7 in /usr/local/lib/python3.10/dist-packages (from roboflow) (2022.12.7)
Requirement already satisfied: chardet==4.0.0 in /usr/local/lib/python3.10/dist-packages (from roboflow) (4.0.0)
Requirement already satisfied: cyclert==0.10.0 in /usr/local/lib/python3.10/dist-packages (from roboflow) (0.10.0)
Requirement already satisfied: idna==2.10 in /usr/local/lib/python3.10/dist-packages (from roboflow) (2.10)
Requirement already satisfied: kiwisolver==1.3.1 in /usr/local/lib/python3.10/dist-packages (from roboflow) (1.4.4)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from roboflow) (3.7.1)
Requirement already satisfied: numpy==1.18.5 in /usr/local/lib/python3.10/dist-packages (from roboflow) (1.22.4)
Requirement already satisfied: opencv-python==4.1.2 in /usr/local/lib/python3.10/dist-packages (from roboflow) (4.7.0.72)
Requirement already satisfied: Pillow==7.1.2 in /usr/local/lib/python3.10/dist-packages (from roboflow) (8.4.0)
Requirement already satisfied: pyparsing==2.4.7 in /usr/local/lib/python3.10/dist-packages (from roboflow) (2.4.7)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.10/dist-packages (from roboflow) (2.8.2)
Requirement already satisfied: python-dotenv in /usr/local/lib/python3.10/dist-packages (from roboflow) (1.0.0)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from roboflow) (2.27.1)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from roboflow) (1.16.0)
Requirement already satisfied: urllib3==1.26.6 in /usr/local/lib/python3.10/dist-packages (from roboflow) (1.26.15)
Requirement already satisfied: wget in /usr/local/lib/python3.10/dist-packages (from roboflow) (3.2)
Requirement already satisfied: tqdm==4.41.0 in /usr/local/lib/python3.10/dist-packages (from roboflow) (4.65.0)
Requirement already satisfied: PyYAML==5.3.1 in /usr/local/lib/python3.10/dist-packages (from roboflow) (6.0)
Requirement already satisfied: requests-toolbelt in /usr/local/lib/python3.10/dist-packages (from roboflow) (1.0.0)
Requirement already satisfied: contourpy==1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->roboflow) (1.0.7)
Requirement already satisfied: fonttools==4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->roboflow) (4.39.3)
Requirement already satisfied: packaging==20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->roboflow) (23.1)
Requirement already satisfied: charset-normalizer==2.0.0 in /usr/local/lib/python3.10/dist-packages (from requests->roboflow) (2.0)
loading Roboflow workspace...
loading Roboflow project...
Downloading Dataset Version Zip in Bee_train-1 to yolov5pytorch: 100% [74072599 / 74072599] bytes
Extracting Dataset Version Zip to Bee_train-1 in yolov5pytorch: 100% [2302/2302] [00:01<00:00, 1961.26it/s]

103_jpg.rf.42af94702d1469def346d4fa4063c7f9.jpg x



데이터셋 주소 입력하고 이미지파일 확인

1. 데이터셋 가져오기 (3)



The screenshot displays the RoboFlow workspace interface. On the left, a file explorer shows a project structure with folders 'Bee_train-1', 'dataset', and 'train'. The 'train' folder contains an 'images' subfolder with numerous JPEG files. The central terminal window shows the output of a command to clone a dataset, listing various requirements already satisfied (e.g., roboflow, certifi, chardet, cyclert, idna, kiwisolver, matplotlib, numpy, opencv-python, Pillow, pyparsing, python-dateutil, python-dotenv, requests, six, urllib3, wget, tqdm, PyYAML, requests-toolbelt, contourpy, fonttools, packaging, charset-normalizer). Below the terminal output, a command prompt shows the execution of `xcld /content` and `git clone https://github.com/ultralytics/yolov5.git`. On the right, a preview of an image file `110_jpg.rf.7ba8765c...` is shown, depicting several bees on a textured surface.

데이터셋 주소 입력하고 이미지파일 확인

2. 옴로 환경 설치

[3] #옴로 클론을 설치하는 코드 디렉토리가 생성됨

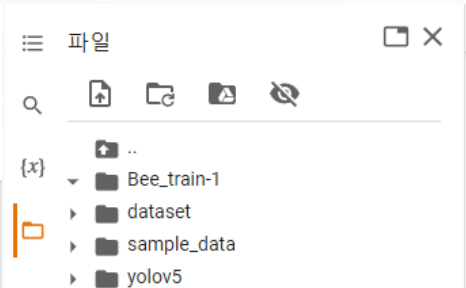
```
%cd /content
!git clone https://github.com/ultralytics/yolov5.git

/content
Cloning into 'yolov5'...
remote: Enumerating objects: 15976, done.
remote: Counting objects: 100% (145/145), done.
remote: Compressing objects: 100% (70/70), done.
remote: Total 15976 (delta 88), reused 116 (delta 75), pack-reused 15831
Receiving objects: 100% (15976/15976), 14.61 MiB | 22.16 MiB/s, done.
Resolving deltas: 100% (10957/10957), done.
```

[4] #옴로5를 위한 패키지 설치

```
%cd /content/yolov5/
!pip install -r requirements.txt

/content/yolov5
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting gitpython>=3.1.30 (from -r requirements.txt (line 5))
  Downloading GitPython-3.1.31-py3-none-any.whl (184 kB)
    184.3/184.3 kB 4.6 MB/s eta 0:00:00
Requirement already satisfied: matplotlib>=3.3 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 6)) (3.7)
Requirement already satisfied: numpy>=1.18.5 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 7)) (1.22)
Requirement already satisfied: opencv-python>=4.1.1 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 8)) (8.4.0)
Requirement already satisfied: Pillow>=7.1.2 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 9)) (8.4.0)
Requirement already satisfied: psutil in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 10)) (5.9.5)
Requirement already satisfied: PyYAML>=5.3.1 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 11)) (6.0)
Requirement already satisfied: requests>=2.23.0 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 12)) (2)
Requirement already satisfied: scipy>=1.4.1 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 13)) (1.10)
Collecting thop>=0.1.1 (from -r requirements.txt (line 14))
  Downloading thop-0.1.1.post2209072238-py3-none-any.whl (15 kB)
Requirement already satisfied: torch>=1.7.0 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 15)) (2.0.1)
Requirement already satisfied: torchvision>=0.8.1 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 16)) (0.11.0)
Requirement already satisfied: tqdm>=4.64.0 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 17)) (4.65)
Collecting ultralytics>=8.0.111 (from -r requirements.txt (line 18))
  Downloading ultralytics-8.0.114-py3-none-any.whl (595 kB)
    595.4/595.4 kB 19.5 MB/s eta 0:00:00
Requirement already satisfied: pandas>=1.1.4 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 27)) (1.5)
Requirement already satisfied: seaborn>=0.11.0 in /usr/local/lib/python3.10/dist-packages (from -r requirements.txt (line 28)) (0
```



옴로 클론 설치, 패키지 설치

3. 이미지저장 리스트생성 및 이미지 개수확인

[6] #data.yaml 파일 확인
%cat /content/dataset/data.yaml

```
names:
- hornet
nc: 1
roboflow:
  license: CC BY 4.0
  project: bee_train
  url: https://universe.roboflow.com/tf2bee/bee\_train/dataset/1
  version: 1
  workspace: tf2bee
test: ../test/images
train: Bee_train-1/train/images
val: Bee_train-1/valid/images
```

[8] #이미지 저장한 리스트 생성 및 이미지개수확인
%cd /
from glob import glob
img_list = glob('content/dataset/train/images/*.jpg')

print(len(img_list))

```
/
1148
```

이미지 리스트 및 개수확인

4. 트레이닝셋과 밸리데이션 셋 나누기 및 텍스트 파일로 데이터셋 저장



The screenshot displays a Jupyter Notebook interface. On the left, a file explorer shows the project structure: a folder named 'Bee_train-1' containing a 'dataset' folder. Inside 'dataset', there is a 'train' folder with files 'README.dataset.txt', 'README.roboflow.txt', 'data.yaml', 'train.txt', and 'val.txt'. There are also 'sample_data' and 'yolov5' folders.

The main area shows two code cells. The first cell, labeled '[10]', contains Python code to split the dataset using `train_test_split` from `sklearn.model_selection`. It splits the `img_list` into `train_img_list` and `val_img_list` with a `test_size` of 0.2 and a `random_state` of 2000. A `print` statement shows the lengths of the resulting lists: 918 for training and 230 for validation.

```
[10] #트레이닝셋과 밸리데이션 나누어 주기
#train_test_split() 트레이닝 셋과 테스트 셋을 나누는 함수
from sklearn.model_selection import train_test_split

train_img_list, val_img_list = train_test_split(img_list, test_size = 0.2, random_state=2000)
#트레이닝 셋은 80퍼센트, 밸리데이션셋은 20퍼센트

print(len(train_img_list), len(val_img_list))
```

The output of the first cell is: 918 230

The second code cell, labeled with a play button icon, contains Python code to save the image IDs from the training and validation lists into text files named 'train.txt' and 'val.txt' respectively, using the `with open` and `f.write` methods.

```
#txt로 데이터셋 저장
with open('/content/dataset/train.txt', 'w') as f:
    f.write('\n'.join(train_img_list) + '\n')

with open('/content/dataset/val.txt', 'w') as f:
    f.write('\n'.join(val_img_list) + '\n')
```

Train, val 값 TXT로 저장

5. 위치변경

✓
1초



#yaml의 위치에 트레이닝셋과 밸리데이션 셋을 넣어주기

```
import yaml
```

```
with open('content/dataset/data.yaml', 'r') as f:  
    data = yaml.safe_load(f)
```

```
print(data)
```

```
data['train'] = '/content/dataset/train.txt'
```

```
data['val'] = '/content/dataset/val.txt'
```

```
with open('/content/dataset/data.yaml', 'w') as f:  
    yaml.dump(data, f)
```

```
print(data)
```

```
on': 1, 'workspace': 'tf2bee'}, 'test': '../test/images', 'train': 'Bee_train-1/train/images', 'val': 'Bee_train-1/valid/images'}  
on': 1, 'workspace': 'tf2bee'}, 'test': '../test/images', 'train': '/content/dataset/train.txt', 'val': '/content/dataset/val.txt'}
```

트레이닝 시키기 위해 위치지정

6. 트레이닝 시키기



#트레이닝 시키기

```
%cd /content/yolov5/
```

```
!python train.py --img 416 --batch 16 --epochs 50 --data /content/dataset/data.yaml --cfg ./models/yolov5m.yaml --weights yolov5m.pt --name bee_yolov5m_results
```

... /content/yolov5

train: weights=yolov5m.pt, cfg=./models/yolov5m.yaml, data=/content/dataset/data.yaml, hyp=data/hyps/hyp.scratch-low.yaml, epochs=50, batch_size=16, imgsz=416, rect=Fal

github: up to date with <https://github.com/ultralytics/yolov5> ✓

YOLOv5 🚀 v7.0-178-ga199480 Python-3.10.11 torch-2.0.1+cu118 CPU

hyperparameters: lr0=0.01, lrf=0.01, momentum=0.937, weight_decay=0.0005, warmup_epochs=3.0, warmup_momentum=0.8, warmup_bias_lr=0.1, box=0.05, cls=0.5, cls_pw=1.0, ot

ClearML: run 'pip install clearml' to automatically track, visualize and remotely train YOLOv5 🚀 in ClearML

Comet: run 'pip install comet_ml' to automatically track and visualize YOLOv5 🚀 runs in Comet

TensorBoard: Start with 'tensorboard --logdir runs/train', view at <http://localhost:6006/>

Overriding model.yaml nc=80 with nc=1

	from	n	params	module	arguments
0		-1	1	models.common.Conv	[3, 48, 6, 2, 2]
1		-1	1	models.common.Conv	[48, 96, 3, 2]
2		-1	2	models.common.C3	[96, 96, 2]
3		-1	1	models.common.Conv	[96, 192, 3, 2]
4		-1	4	models.common.C3	[192, 192, 4]
5		-1	1	models.common.Conv	[192, 384, 3, 2]
6		-1	6	models.common.C3	[384, 384, 6]
7		-1	1	models.common.Conv	[384, 768, 3, 2]
8		-1	2	models.common.C3	[768, 768, 2]
9		-1	1	models.common.SPPF	[768, 768, 5]
10		-1	1	models.common.Conv	[768, 384, 1, 1]
11		-1	1	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
12		[-1, 6]	1	models.common.Concat	[1]
13		-1	2	models.common.C3	[768, 384, 2, False]
14		-1	1	models.common.Conv	[384, 192, 1, 1]
15		-1	1	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
16		[-1, 4]	1	models.common.Concat	[1]
17		-1	2	models.common.C3	[384, 192, 2, False]
18		-1	1	models.common.Conv	[192, 192, 3, 2]
19		[-1, 14]	1	models.common.Concat	[1]
20		-1	2	models.common.C3	[384, 384, 2, False]
21		-1	1	models.common.Conv	[384, 384, 3, 2]
22		[-1, 10]	1	models.common.Concat	[1]

6. 트레이닝 시키기

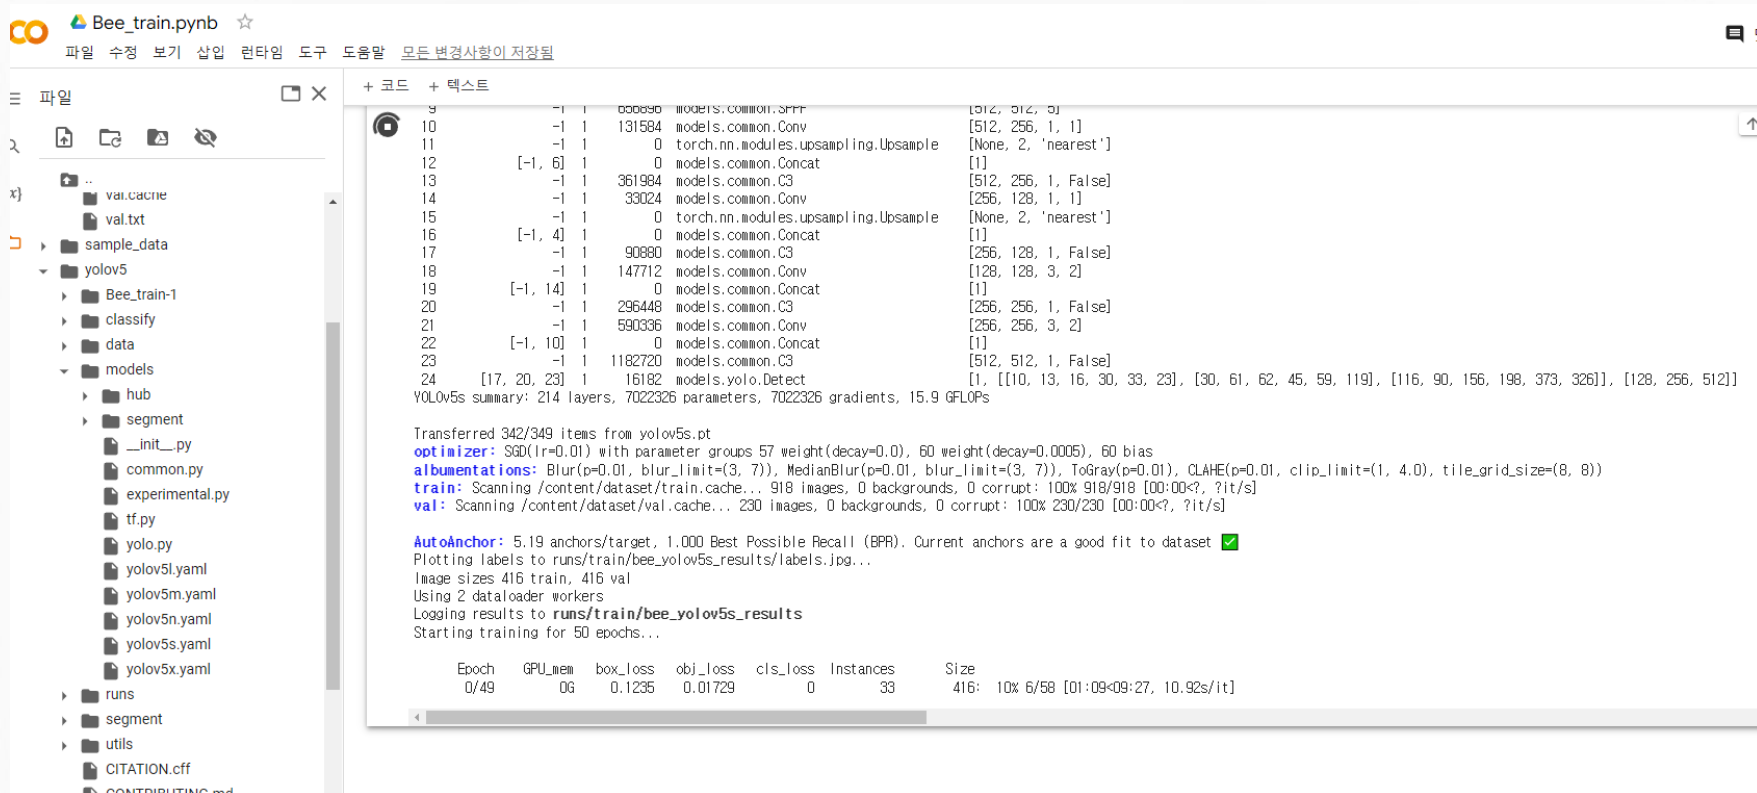
```
6      -1 6 2512896 models.common.C3 [384, 384, 6]
7      -1 1 2655744 models.common.Conv [384, 768, 3, 2]
8      -1 2 4134912 models.common.C3 [768, 768, 2]
9      -1 1 1476864 models.common.SPPF [768, 768, 5]
10     -1 1 295680 models.common.Conv [768, 384, 1, 1]
11     -1 1 0 torch.nn.modules.upsampling.Upsample [None, 2, 'nearest']
12     [-1, 6] 1 0 models.common.Concat [1]
13     -1 2 1182720 models.common.C3 [768, 384, 2, False]
14     -1 1 74112 models.common.Conv [384, 192, 1, 1]
15     -1 1 0 torch.nn.modules.upsampling.Upsample [None, 2, 'nearest']
16     [-1, 4] 1 0 models.common.Concat [1]
17     -1 2 296448 models.common.C3 [384, 192, 2, False]
18     -1 1 332160 models.common.Conv [192, 192, 3, 2]
19     [-1, 14] 1 0 models.common.Concat [1]
20     -1 2 1035264 models.common.C3 [384, 384, 2, False]
21     -1 1 1327872 models.common.Conv [384, 384, 3, 2]
22     [-1, 10] 1 0 models.common.Concat [1]
23     -1 2 4134912 models.common.C3 [768, 768, 2, False]
24     [17, 20, 23] 1 24246 models.yolo.Detect [1, [[10, 13, 16, 30, 33, 23], [30, 61, 62, 45, 59, 119], [116, 90, 156, 198, 373, 326]], [192, 384, 768]]
YOLOv5m summary: 291 layers, 20871318 parameters, 20871318 gradients, 48.2 GFLOPs

Transferred 474/481 items from yolov5m.pt
optimizer: SGD(lr=0.01) with parameter groups 79 weight(decay=0.0), 82 weight(decay=0.0005), 82 bias
augmentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01), CLAHE(p=0.01, clip_limit=(1, 4.0), tile_grid_size=(8, 8))
train: Scanning /content/dataset/train... 918 images, 0 backgrounds, 0 corrupt: 100% 918/918 [00:03<00:00, 271.83it/s]
train: New cache created: /content/dataset/train.cache
val: Scanning /content/dataset/val... 230 images, 0 backgrounds, 0 corrupt: 100% 230/230 [00:00<00:00, 269.37it/s]
val: New cache created: /content/dataset/val.cache

AutoAnchor: 5.19 anchors/target, 1.000 Best Possible Recall (BPR). Current anchors are a good fit to dataset ✓
Plotting labels to runs/train/bee_yolov5m_results3/labels.jpg...
Image sizes 416 train, 416 val
Using 2 dataloader workers
Logging results to runs/train/bee_yolov5m_results3
Starting training for 50 epochs...

Epoch  GPU_mem  box_loss  obj_loss  cls_loss  Instances  Size
0/49    0G       0.1144   0.01823   0         38         416: 2% 1/58 [00:37<35:54, 37.80s/it]
```

6. 트레이닝 시키기 모델 s로 변경, 학습 수 10번으로 변경



```
9      -1 1 00090 models.common.SPP [512, 512, 5]
10     -1 1 131584 models.common.Conv [512, 256, 1, 1]
11     -1 1 0 torch.nn.modules.upsampling.Upsample [None, 2, 'nearest']
12     [-1, 6] 1 0 models.common.Concat [1]
13     -1 1 361984 models.common.C3 [512, 256, 1, False]
14     -1 1 33024 models.common.Conv [256, 128, 1, 1]
15     -1 1 0 torch.nn.modules.upsampling.Upsample [None, 2, 'nearest']
16     [-1, 4] 1 0 models.common.Concat [1]
17     -1 1 90880 models.common.C3 [256, 128, 1, False]
18     -1 1 147712 models.common.Conv [128, 128, 3, 2]
19     [-1, 14] 1 0 models.common.Concat [1]
20     -1 1 296448 models.common.C3 [256, 256, 1, False]
21     -1 1 590336 models.common.Conv [256, 256, 3, 2]
22     [-1, 10] 1 0 models.common.Concat [1]
23     -1 1 1182720 models.common.C3 [512, 512, 1, False]
24     [17, 20, 23] 1 16182 models.yolo.Detect [1, [[10, 13, 16, 30, 33, 23], [30, 61, 62, 45, 59, 119], [116, 90, 156, 198, 373, 326]], [128, 256, 512]]
YOLOv5s summary: 214 layers, 7022326 parameters, 7022326 gradients, 15.9 GFLOPs

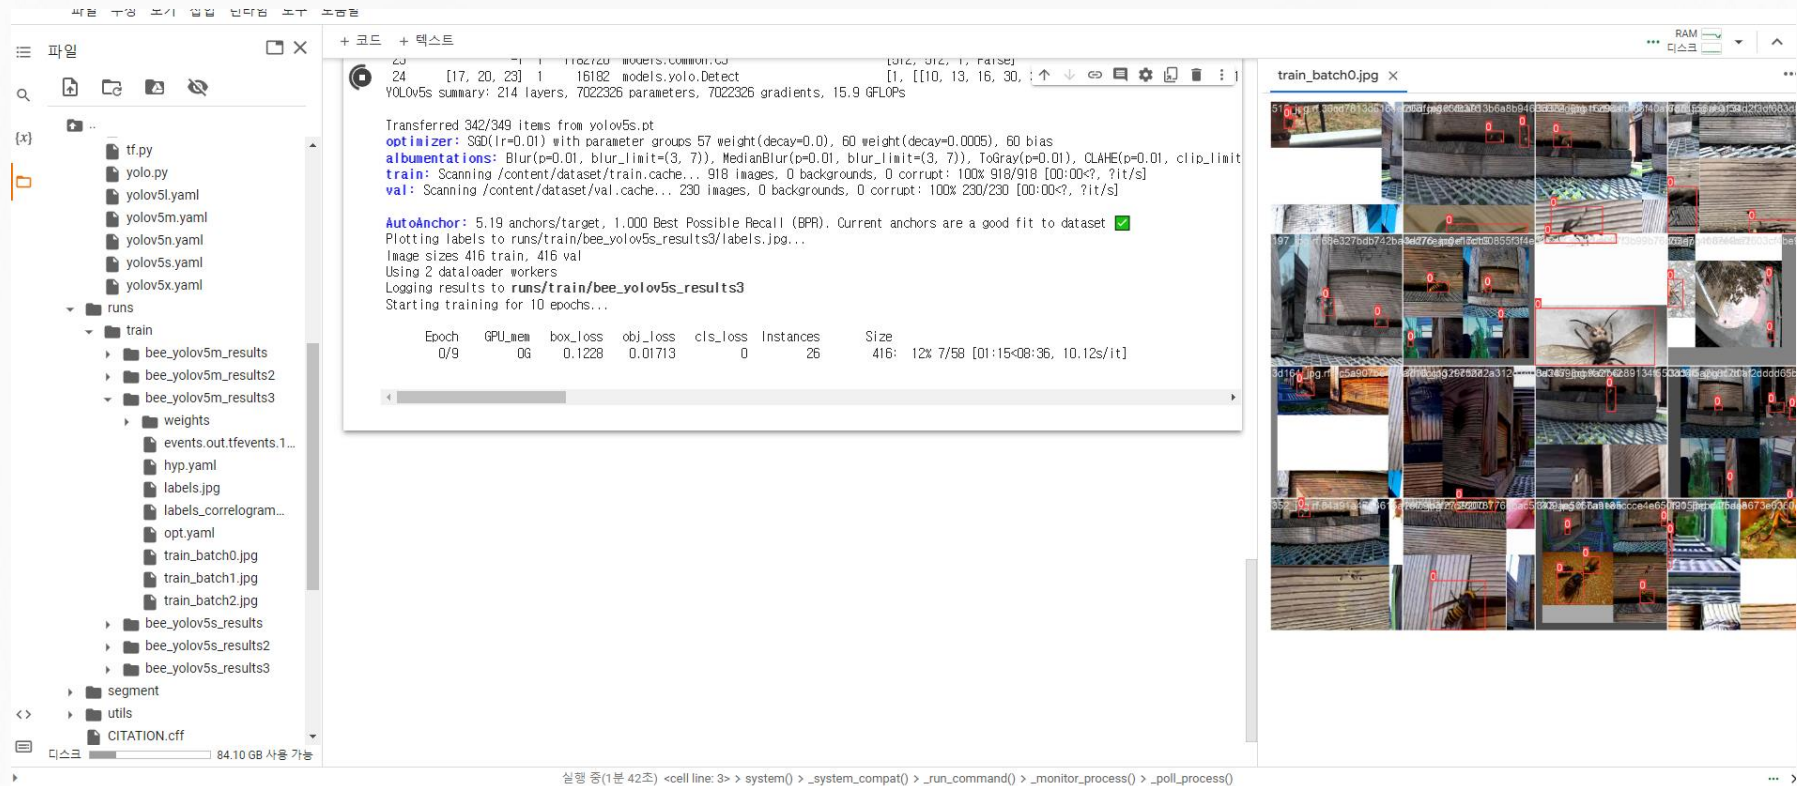
Transferred 342/349 items from yolov5s.pt
optimizer: SGD(lr=0.01) with parameter groups 57 weight(decay=0.0), 60 weight(decay=0.0005), 60 bias
augmentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01), CLAHE(p=0.01, clip_limit=(1, 4.0), tile_grid_size=(8, 8))
train: Scanning /content/dataset/train.cache... 918 images, 0 backgrounds, 0 corrupt: 100% 918/918 [00:00<?, ?it/s]
val: Scanning /content/dataset/val.cache... 230 images, 0 backgrounds, 0 corrupt: 100% 230/230 [00:00<?, ?it/s]

AutoAnchor: 5.19 anchors/target, 1.000 Best Possible Recall (BPR). Current anchors are a good fit to dataset ✓
Plotting labels to runs/train/bee_yolov5s_results/labels.jpg...
Image sizes 416 train, 416 val
Using 2 dataloader workers
Logging results to runs/train/bee_yolov5s_results
Starting training for 50 epochs...

Epoch  GPU_mem  box_loss  obj_loss  cls_loss  Instances  Size
0/49      OG    0.1235   0.01729    0         33        416: 10% 6/58 [01:09<09:27, 10.92s/it]
```

속도향상을 위해 코랩 GPU 활용 및 모델사이즈 s로 변경, 학습 수 10번으로 변경

6. 트레이닝 모델 배치활동 확인



The screenshot displays a Jupyter Notebook interface with a file explorer on the left and a code cell in the center. The file explorer shows a directory structure for a YOLOv5 project, including files like `tf.py`, `yolo.py`, `yolov5l.yaml`, `yolov5m.yaml`, `yolov5n.yaml`, `yolov5s.yaml`, and `yolov5x.yaml`. The code cell shows the output of a YOLOv5 training session, including a summary of the model, training progress, and a table of training metrics.

YOLOv5s summary: 214 layers, 7022326 parameters, 7022326 gradients, 15.9 GFLOPs

Transferred 342/349 items from yolov5s.pt

optimizer: SGD(lr=0.01) with parameter groups 57 weight(decay=0.0), 60 weight(decay=0.0005), 60 bias

augmentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01), CLAHE(p=0.01, clip_limit=2, tile_grid_size=(8, 8))

train: Scanning /content/dataset/train.cache... 918 images, 0 backgrounds, 0 corrupt: 100% 918/918 [00:00<?, ?it/s]

val: Scanning /content/dataset/val.cache... 230 images, 0 backgrounds, 0 corrupt: 100% 230/230 [00:00<?, ?it/s]

AutoAnchor: 5.19 anchors/target, 1.000 Best Possible Recall (BPR). Current anchors are a good fit to dataset ✓

Plotting labels to runs/train/bee_yolov5s_results3/labels.jpg...

Image sizes 416 train, 416 val

Using 2 dataloader workers

Logging results to runs/train/bee_yolov5s_results3

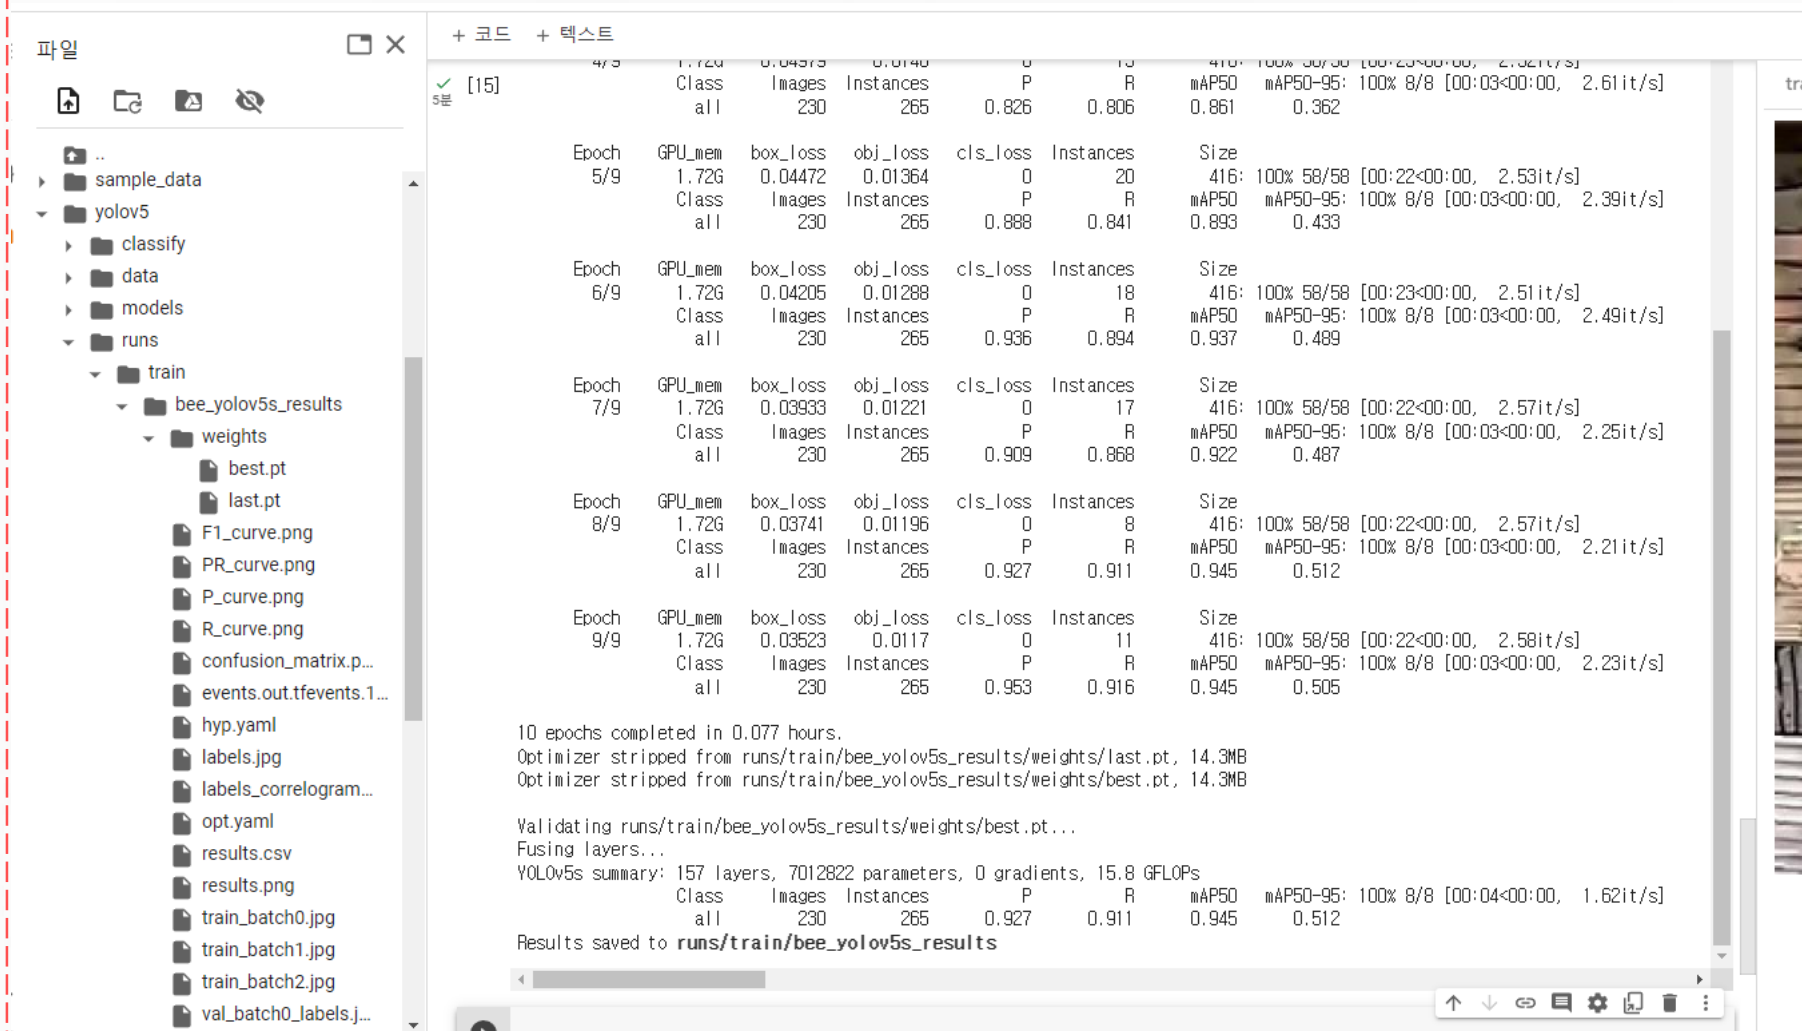
Starting training for 10 epochs...

Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size
0/9	0G	0.1228	0.01713	0	26	416x 12x 7/58 [01:15<08:36, 10.12s/it]

The right side of the image shows a grid of training images with bounding boxes, labeled `train_batch0.jpg`.

실행 중(1분 42초) <cell line: 3> > system() > _system_compat() > _run_command() > _monitor_process() > _poll_process()

7. 트레이닝 완료하고 확인하기



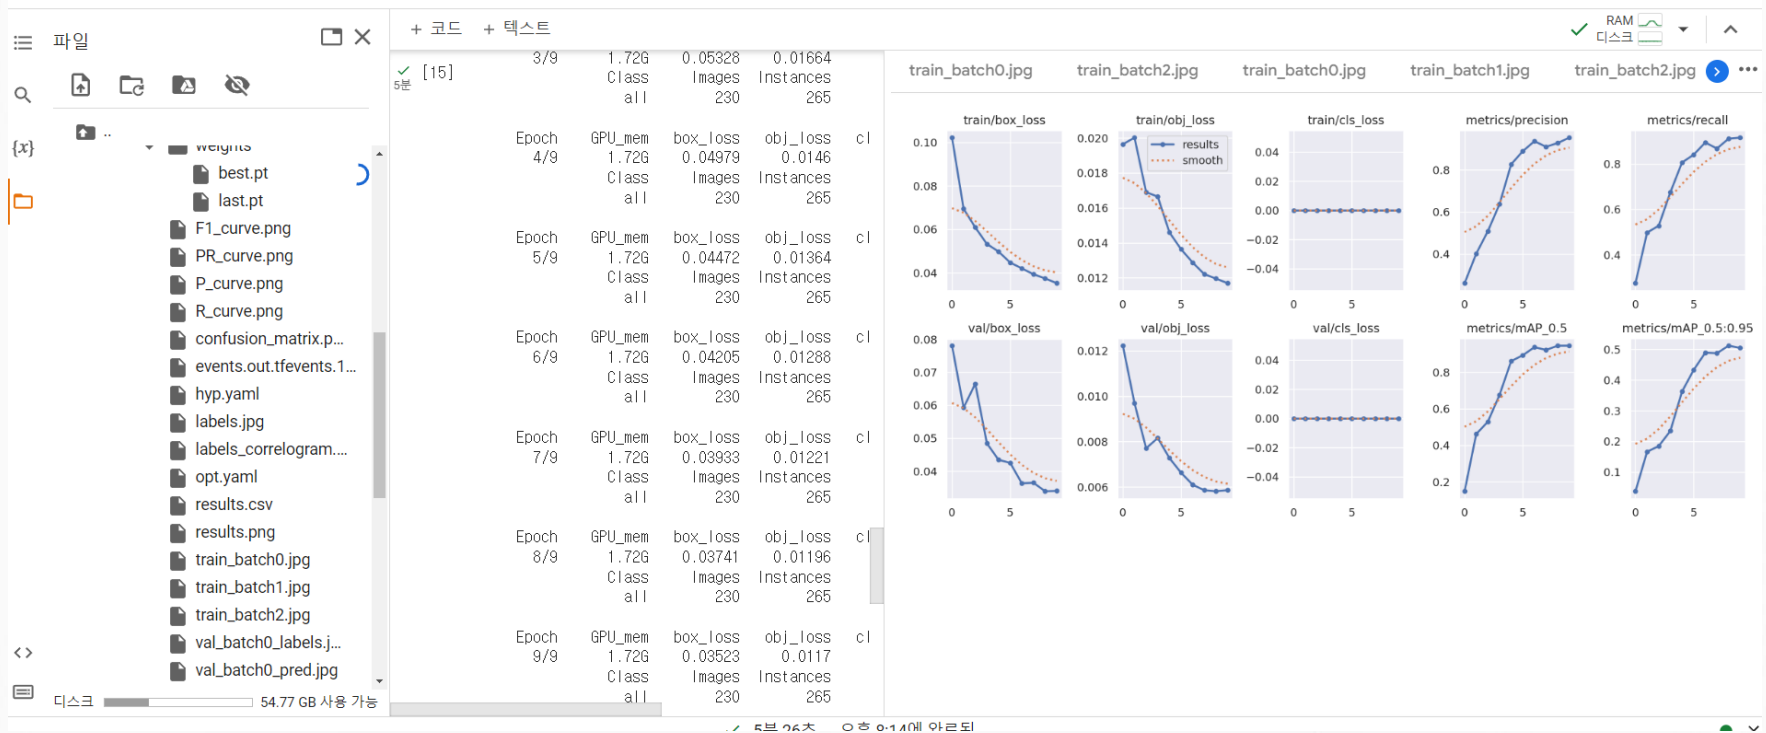
The image shows a file explorer on the left and a terminal window on the right. The file explorer displays the directory structure of a YOLOv5 training run, including sample_data, yolov5, classify, data, models, runs, and train. The train directory contains a subdirectory named bee_yolov5s_results, which includes weights, best.pt, last.pt, and various performance metrics (F1_curve.png, PR_curve.png, P_curve.png, R_curve.png, confusion_matrix.p..., events.out.tfevents.1..., hyp.yaml, labels.jpg, labels_correlogram..., opt.yaml, results.csv, results.png, train_batch0.jpg, train_batch1.jpg, train_batch2.jpg, val_batch0_labels.j...).

The terminal window shows the output of the training process. It displays the progress of the training, including the number of epochs completed (10 epochs), the time taken (0.077 hours), and the results of the validation process. The terminal output includes the following text:

```
10 epochs completed in 0.077 hours.  
Optimizer stripped from runs/train/bee_yolov5s_results/weights/last.pt, 14.3MB  
Optimizer stripped from runs/train/bee_yolov5s_results/weights/best.pt, 14.3MB  
  
Validating runs/train/bee_yolov5s_results/weights/best.pt...  
Fusing layers...  
YOLOv5s summary: 157 layers, 7012822 parameters, 0 gradients, 15.8 GFLOPs  
Class Images Instances P R mAP50 mAP50-95 100% 8/8 [00:04<00:00, 1.62it/s]  
all 230 265 0.927 0.911 0.945 0.512  
  
Results saved to runs/train/bee_yolov5s_results
```

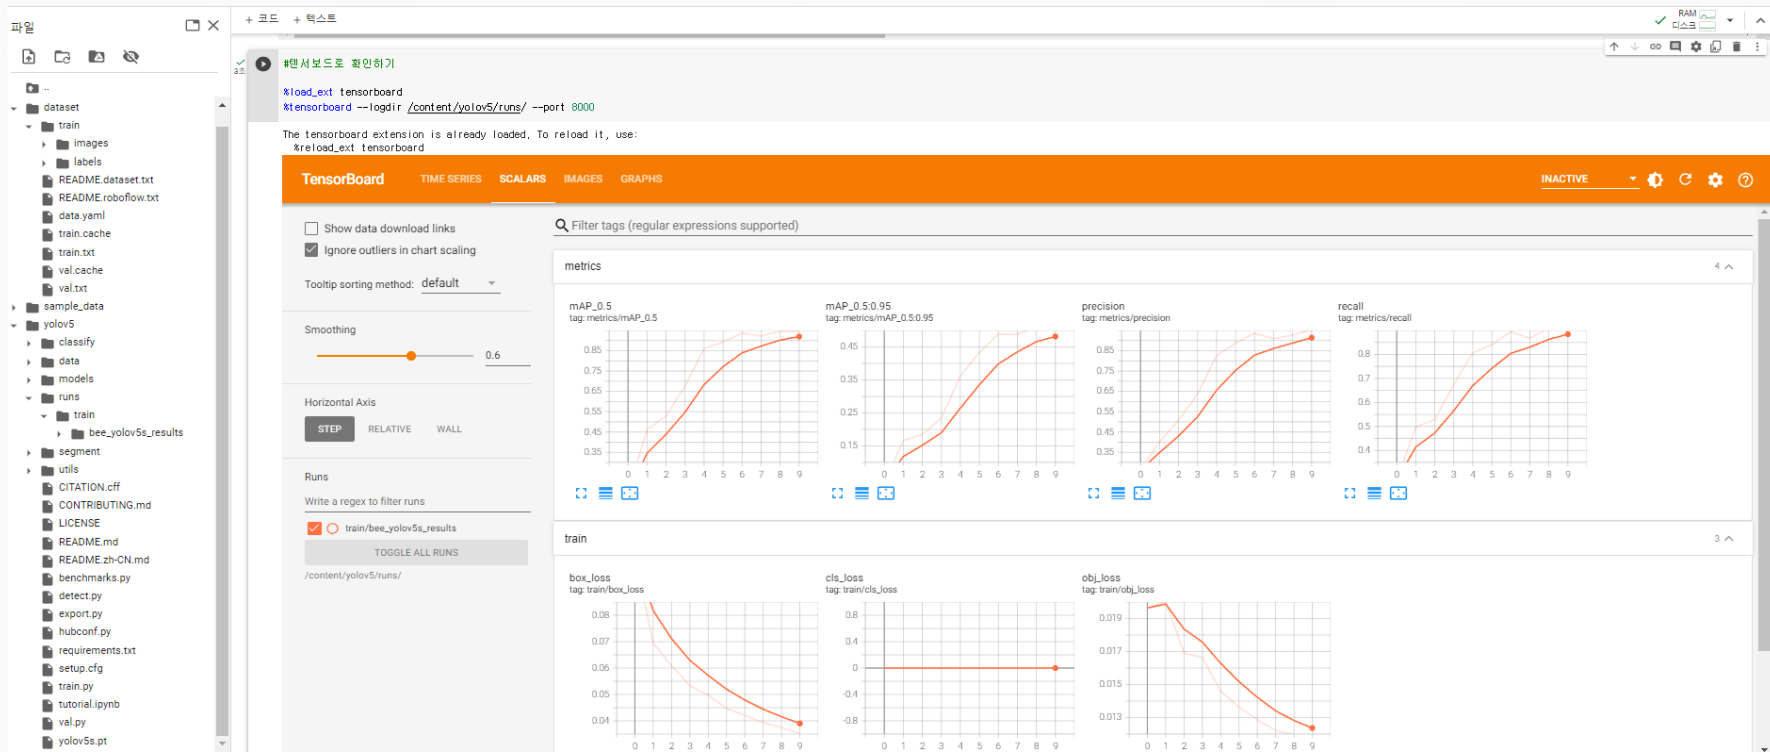
Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size	mAP50	mAP50-95	100%	8/8	[00:03<00:00, 2.61it/s]
4/9	1.72G	0.04375	0.0140	0	18	416	0.861	0.362	100%	8/8	[00:03<00:00, 2.61it/s]
5/9	1.72G	0.04472	0.01364	0	20	416	0.893	0.433	100%	8/8	[00:03<00:00, 2.39it/s]
6/9	1.72G	0.04205	0.01288	0	18	416	0.937	0.489	100%	8/8	[00:03<00:00, 2.49it/s]
7/9	1.72G	0.03933	0.01221	0	17	416	0.922	0.487	100%	8/8	[00:03<00:00, 2.25it/s]
8/9	1.72G	0.03741	0.01196	0	8	416	0.945	0.512	100%	8/8	[00:03<00:00, 2.21it/s]
9/9	1.72G	0.03523	0.0117	0	11	416	0.945	0.505	100%	8/8	[00:03<00:00, 2.23it/s]

7. 트레이닝 완료하고 확인하기



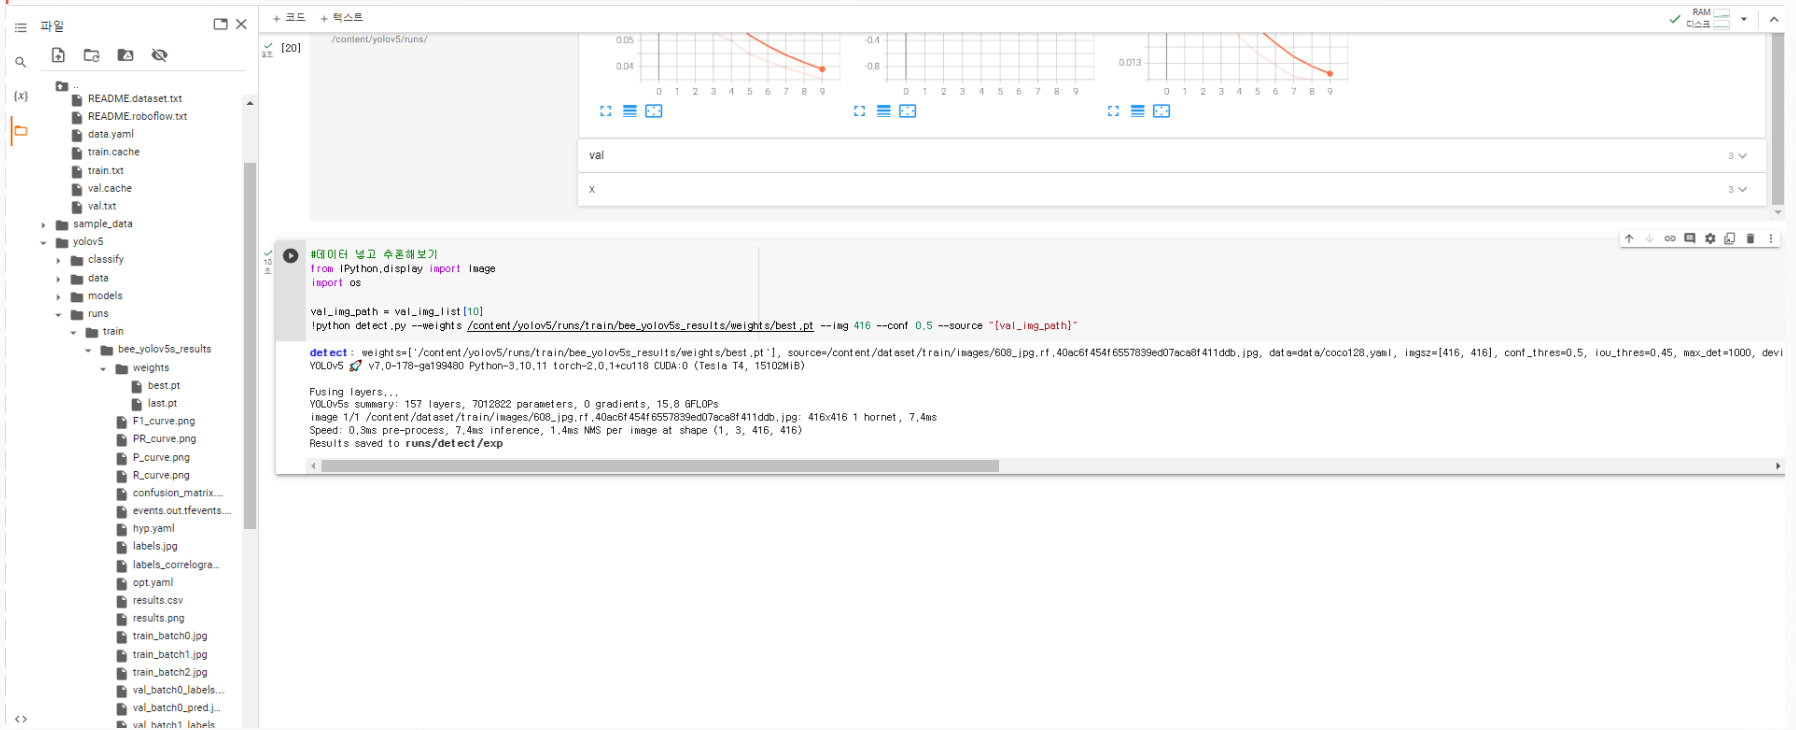
결과값 확인

7. 트레이닝 완료하고 확인하기



텐서플로우로 결과값 확인

8. 이미지 삽입 후 테스트



데이터 넣고 추론하기
runs/detect/exp로 저장됨

8. 이미지 삽입 후 테스트

Bee_train.pynb ☆

파일 수정 보기 삽입 런타임 도구 도움말 모든 변경사항이 저장됨


파일

- (x) ..
- Bee_train-1
 - dataset
 - train
 - images
 - labels
 - README.dataset.txt
 - README.roboflow.txt
 - data.yaml
 - train.cache
 - train.txt
 - val.cache
 - val.txt
 - sample_data
 - yolov5
 - classify
 - data
 - models
 - runs
 - detect
 - exp
 - 608_jpg.rf.40ac6f45...

코드

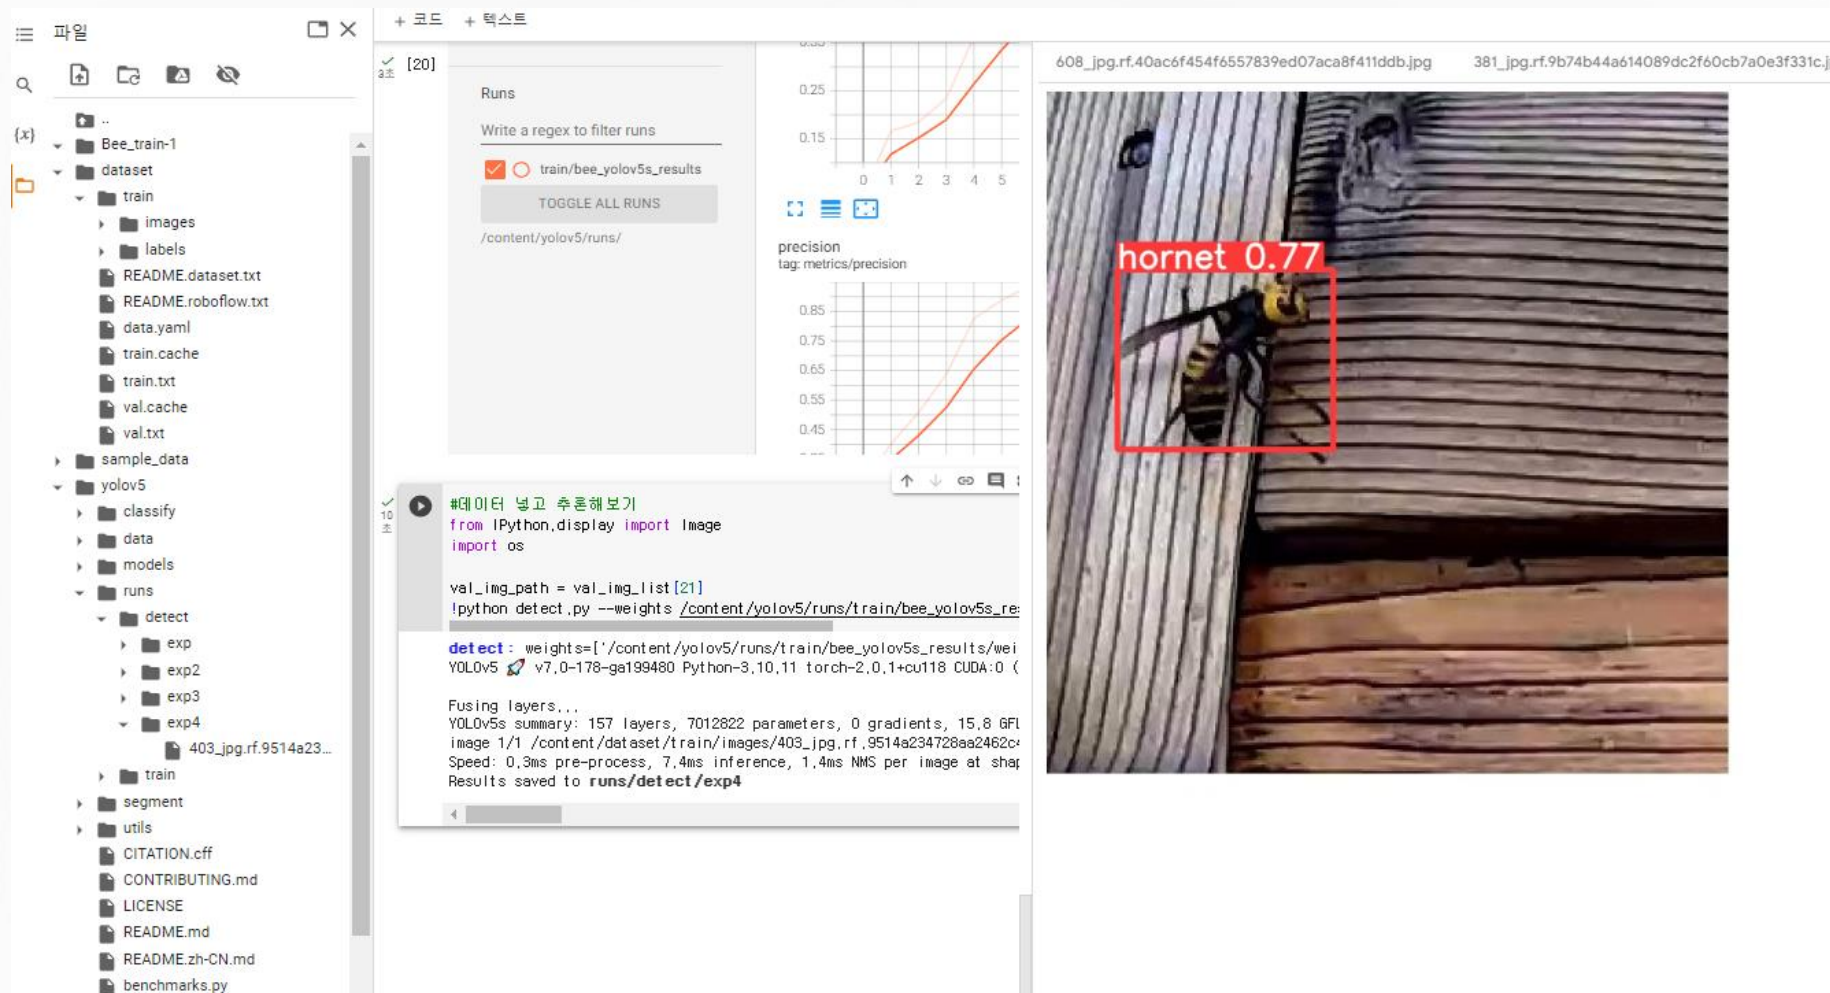
```
[20] TOGGLE ALL RUNS  
/content/yolov5/runs/  
  
#데이터 넣고 추론해 보기  
from IPython.display import Image  
import os  
  
val_img_path = val_img_list[10]  
!python detect.py --weights /content.  
  
detect: weights=[ '/content/yolov5/r  
YOLOv5 v7.0-178-ga199480 Python-3  
  
Fusing layers...  
YOLOv5s summary: 157 layers, 7012822  
Image 1/1 /content/dataset/train/imag  
Speed: 0.3ms pre-process, 7.4ms infer  
Results saved to runs/detect/exp
```

608_jpg.rf.40ac6f454f6557839ed07aca8f411ddb.jpg ×



Val[10]번에 있는 이미지를 가져와서 넣어보니
잘 잡아내는 모습을 확인

8. 이미지 삽입 후 테스트



The screenshot displays a Jupyter Notebook environment. On the left, a file explorer shows a directory structure for a YOLOv5 project, including 'Bee_train-1', 'dataset', 'train', 'sample_data', and 'yolov5'. The main area shows a code cell with the following content:

```
[20]
Runs
Write a regex to filter runs
☒ ☐ train/bee_yolov5s_results
TOGGLE ALL RUNS
/content/yolov5/runs/

#데이터 넣고 추론해보기
from IPython.display import Image
import os

val_img_path = val_img_list[21]
!python detect.py --weights /content/yolov5/runs/train/bee_yolov5s_re

detect: weights=['/content/yolov5/runs/train/bee_yolov5s_results/wei
YOLOv5 v7.0-178-ga199480 Python-3.10.11 torch-2.0.1+cu118 CUDA:0 (
Fusing layers...
YOLOv5s summary: 157 layers, 7012822 parameters, 0 gradients, 15.8 GFLOPs
image 1/1 /content/dataset/train/images/403.jpg.rf.9514a234728aa2462c
Speed: 0.3ms pre-process, 7.4ms inference, 1.4ms NMS per image at shap
Results saved to runs/detect/exp4
```

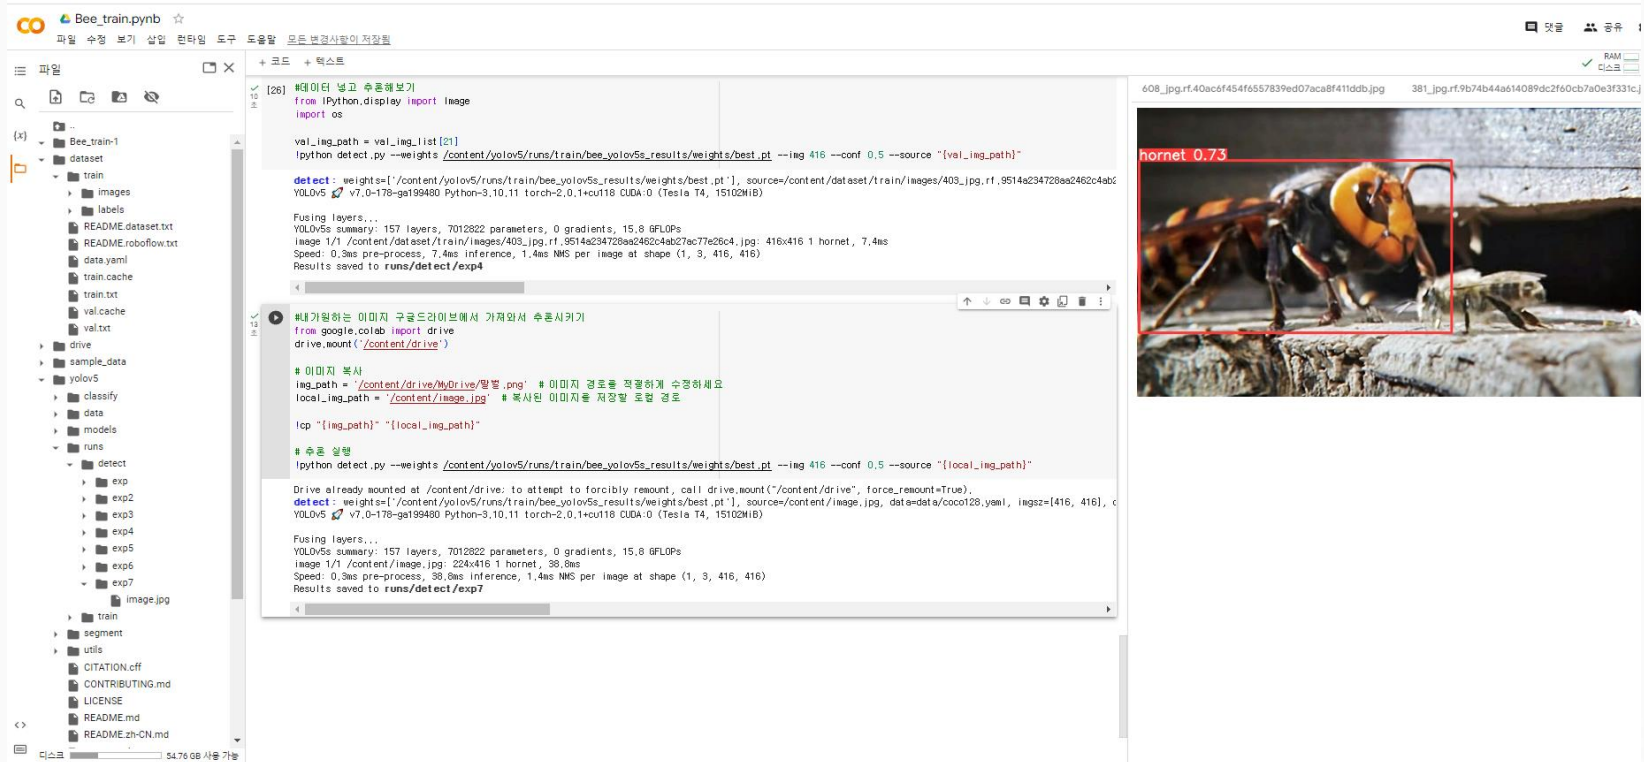
Below the code cell, a terminal window shows the execution output. To the right of the code cell, two line graphs are visible: the top one shows 'precision' (y-axis 0.15 to 0.25) vs an unlabeled x-axis (0 to 5), and the bottom one shows 'tag: metrics/precision' (y-axis 0.45 to 0.85) vs an unlabeled x-axis (0 to 5). On the far right, a detection result image is shown, featuring a hornet on a wooden surface, enclosed in a red bounding box with the label 'hornet 0.77'.

Val[21]번에 있는 이미지를 가져와서 넣어보니
잘 잡아내는 모습을 확인

9. 원하는 이미지 업로드 후 테스트



9. 원하는 이미지 업로드 후 테스트



```
[26] #데이터 넣고 수행해보기
from IPython.display import Image
import os

val_img_path = val_img_list[21]
!python detect.py --weights /content/yolov5/runs/train/bee_yolov5s_results/weights/best.pt --img 416 --conf 0.5 --source "[val_img_path]"

YOLOv5 v7.0-178-ga199480 Python-3.10.11 torch-2.0.1+cu118 CUDA-0 (Tesla T4, 15102MiB)

Fusing layers...
YOLOv5s summary: 157 layers, 7012822 parameters, 0 gradients, 15.8 GFLOPs
image 1/1 /content/dataset/train/images/403.jpg.rf.9514a234728a2462c4ab27ec77e26c4.jpg: 416x416 1 hornet, 7.4ms
Speed: 0.3ms pre-process, 7.4ms inference, 1.4ms NMS per image at shape (1, 3, 416, 416)
Results saved to runs/detect/exp4

#내가 원하는 이미지 구글드라이브에서 가져와서 수행시키기
from google.colab import drive
drive.mount('/content/drive')

# 이미지 복사
img_path = '/content/drive/MyDrive/말벌.png' # 이미지 경로를 적절하게 수정하세요
local_img_path = '/content/image.jpg' # 복사된 이미지를 저장할 로컬 경로

!cp "{img_path}" "{local_img_path}"

# 수행 실행
!python detect.py --weights /content/yolov5/runs/train/bee_yolov5s_results/weights/best.pt --img 416 --conf 0.5 --source "[local_img_path]"

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount('/content/drive', force_remount=True).
detect: weights=[/content/yolov5/runs/train/bee_yolov5s_results/weights/best.pt], source=/content/image.jpg, data=/content/coco128.yaml, imgsz=[416, 416], c
YOLOv5 v7.0-178-ga199480 Python-3.10.11 torch-2.0.1+cu118 CUDA-0 (Tesla T4, 15102MiB)

Fusing layers...
YOLOv5s summary: 157 layers, 7012822 parameters, 0 gradients, 15.8 GFLOPs
image 1/1 /content/image.jpg: 224x416 1 hornet, 38.8ms
Speed: 0.3ms pre-process, 38.8ms inference, 1.4ms NMS per image at shape (1, 3, 416, 416)
Results saved to runs/detect/exp7
```

608.jpg.rf.40a6cf454f6557839ed07aca8f411ddb.jpg 381.jpg.rf.9b74b44a614089dc2f60cb7a0e3f331c.jpg

hornet 0.73

구글드라이브에 업로드한 나의 이미지로 추론하기
정상적 인식 확인

감사합니다